


Full Training-to-Inference Architecture for Few-Shot 3D Talking Head Generation

Input: video frame sequence + audio • Output: saved weights / checkpoint that can be driven by another audio



Practical use: train once on a target identity -> save checkpoint weights -> reuse the weights with any new audio at inference time.

If you want a simpler baseline, skip Stage A and train only Stage B; if you want a stronger paper, keep Stage A + Stage B.